

LUIS BORGES

Irving, TX | +1 469.279.9909 | luiborgem@gmail.com | linkedin.com/in/luis-borges-bsee-mba-16369a3 |
luiberge23.github.io

Technical Program Manager — AI Infrastructure, Hardware Platforms & AI/ML Delivery

SUMMARY

Technical Program Manager with 15+ years delivering AI infrastructure and hardware programs from concept through hyperscale production, plus 6 years in product management. At Meta, drove AI/GPU infrastructure programs: GB200 (NVLink-connected GPU racks) NPI through EVT/DVT/PVT/MP gates, a 16→256 GPU research-cluster build-out, and dark-fiber data center connectivity. Shipped an AI-powered (RAG/LLM/VLM) analytics platform adopted by 25+ teams in six months, and scaled a validation platform from 2 to 36+ people across 23 testbeds. Former Nokia LTE/5G Product Manager owning a \$50M–\$100M portfolio on the Mobile Broadband Executive Committee. Deep in OEM/ODM and vendor governance, quantitative risk-driven program execution, and matrixed cross-functional leadership across hardware, software, data center, network, and capacity planning teams. MBA and BSEE.

CORE COMPETENCIES

AI Infrastructure & Hardware Platforms: NVLink rack-scale systems (GB200) • NPI gate management (EVT/DVT/PVT/MP) • large-scale AI cluster build-out (16→256 GPUs) • liquid-cooled rack integration and CDU/thermal management • ODM/vendor governance and SOW negotiation • data center deployment and dark-fiber connectivity

AI/ML Product Delivery: shipped LLM/VLM-powered analytics platform to 25+ teams • RAG/LLM/VLM workflows • AI-accelerated validation (analysis from days to minutes) • product roadmaps, strategy, and portfolio ownership (\$50M–\$100M)

Program Leadership at Scale: large-scale cross-organizational programs • quantitative risk registers, dependency management, and escalation • executive reporting and stakeholder alignment • team scaling (2→36+) and coaching • process improvement across functions

Technical Toolkit: Python • Bash • SQL • Jira • Git • Wireshark • Kubernetes • OpenStack / VMware • TCP/IP, QUIC/HTTP-3 • 3GPP / LTE / 5G NR

PROFESSIONAL EXPERIENCE

META — Menlo Park, CA

Technical Program Manager, AI Infrastructure & Hardware Programs | June 2021 – July 2026

AI Hardware Platform NPI (GB200 / NVLink GPU Racks)

- Owned production gate decisions (stop / conditional / go) through EVT, DVT, PVT, and MP; halted a 72-hour burn-in when trapped air in the CDU liquid loop halved flow, turning the fix into a mandatory de-air procedure for every rack.
- Drove NPI diagnostic accuracy to a standing numeric exit criterion via parallel tooling and BMC-firmware tracks across hardware, firmware, and ODM teams; managed strategic vendor engagement and delivery with Keysight and Broadcom.
- Ran program risk registers with quantified schedule exposure across hardware, firmware, and vendor tracks; drove mitigation plans, dependency management, and executive escalation to protect program dates.

Large-Scale AI Cluster & Data Center Delivery

- Scaled AI research infrastructure 12× (16→256 GPUs): put the custom Athena cluster network on the critical path and drove fabric build-out validation, coordinating with capacity planning, data center, and network

infrastructure teams on scale-out sequencing.

- Recovered a de-staffed P0 network path for the Codec Avatars (Reality Labs) research cluster — re-engaged vendors, backfilled engineers, and re-planned POP-to-data-center connectivity, compressing a 9-month plan to 4 months and closing a VP-level escalation.
- Stood up secure remote access into the locked-down Project Aria research cluster from scratch — device certificates, connectivity validation, and a live demo at the RL Symposium — aligning security, network, and research stakeholders with no pre-existing approval process.

AI/ML Product Delivery & Platform Adoption

- Shipped an AI-powered analytics platform (RAG/LLM/VLM workflows) adopted by 25+ teams in six months, cutting hardware validation analysis from days to minutes; defined scope, requirements, and rollout strategy with engineering partners and drove org-wide adoption.
- Scaled a quality-of-experience validation platform for Meta's Family of Apps from 2 to 36+ people and 700+ devices across 23 testbeds (2.3M+ automated jobs/year); secured \$4M+ in multi-year funding and raised uptime from 80% to 96%+.

VERIZON — Southlake, TX

Senior Principal Systems Engineer | July 2020 – June 2021

- Integrated 5G SA core stacks in Kubernetes and delivered Dynamic Spectrum Sharing simulations for carrier-class, cloud-native deployment validation.

NOKIA — Irving, TX

LTE & 5G Product Manager → Services Product Manager | January 2016 – July 2020

- Owned the IoT, 5G, and edge product portfolio across North America and Latin America; drove \$50M–\$100M annual RAN/Core programs through carrier deployment and served on the Mobile Broadband Executive Committee.

NOKIA — Irving, TX

Senior Technical Program Manager / Network Engineering | January 2005 – December 2015

- Led nationwide network deployments and EPC architecture; designed business continuity for mission-critical Packet Core and virtualized multi-tenant validation environments (OpenStack, VMware).

EDUCATION & CERTIFICATIONS

MBA — University of Dallas, Irving, TX

BSEE, Electrical Engineering — Universidad Central de Venezuela

CS50 Python (Harvard/edX, 2023) | 5G NR (Nokia, 2017) | Certified Network Associate Cloud (Nokia, 2016) | OpenStack Administrator (Mirantis, 2015) | Mobile Packet Core / EPC (Nokia, 2012)